

Aman Chauhan

 linkedin.com/aman-chauhan

 +919131118509

 Bengaluru, India

 amanchauhan.iitian@gmail.com

EXPERIENCE

Meesho

Bengaluru, India

Software Development Engineer - III

Jan 2026 - Current

- Designed and scaled an org-wide **Text-to-SQL platform** on internal notebook infrastructure, enabling automatic schema discovery, metadata indexing, and standardized query repositories across platforms.
- Built an **AI-powered** Text-to-SQL system enabling **natural language-based SQL generation** and execution at a scale of 5K queries per day, achieving **70%+ query accuracy** and 75% adherence to SQL best practices.

Software Development Engineer - II

Jan 2024 - Dec 2025

- Engineered a **scalable, distributed notebook platform on Spark + Kubernetes** supporting **500+ analysts and ML engineers**, enabling multi-language workloads (Python, SQL, Scala) with isolated runtime environments. Architected **JupyterHub + Enterprise Gateway** for multi-tenant kernel execution with SSO and integrations with **Cursor and VS Code**.
- Developed a **Golang-based control plane** to orchestrate Spark clusters on Kubernetes, handling autoscaling, kernel lifecycle management, idle termination, dependency isolation, and DBFS-like file synchronization with **Hive, Delta Lake, and GCP connectors**. Integrated Gluten (Velox) for faster Spark performance.
- Built an internal notebook platform with **Databricks** like capabilities on Spark + Kubernetes, including compute provisioning, job scheduling, notebook lifecycle management, and workflow automation. Reduced cost per order by 0.1 Rs., resulting in **Rs.5 crore in annual savings**, with adoption across 500+ users.
- Enhanced query execution across **Zeppelin, Models, and Superset** by implementing a **fail-fast mechanism** with strict query validation, including partition usage enforcement.
- Leveraged **SQLGlott** for query lineage analysis, ensuring correct partition usage, restricting access to sensitive tables, and reducing **infrastructure costs by 30%**.
- Transitioned from cron-based scheduling to an **event-driven, lineage-aware scheduling** system, enabling tasks to trigger based on upstream dependencies. Improved scheduling, lowering **SLA delays by 60%**, eliminated stale data issues, and improved data accuracy and resource utilization.

Software Development Engineer - I

July 2022 - Dec 2023

- Optimized **ETL server costs by 50% (Rs.500,000/day to Rs.250,000/day)** via resource optimization and cost-effective instance migration. Tuned Spark configs to cut excess workers and improve SLA adherence.
- Implemented a comprehensive **observability framework** using Telegraf and VictoriaMetrics with real-time Grafana dashboards, and designed an **Anomaly Detection system** (standard deviation, IQR, z-scores) for early detection of data inconsistencies and performance issues.
- Led migration of critical infrastructure components, including Databricks workspace, Airflow EC2 to EKS, and Kafka Connect MSK to Confluent Kafka. Delivered a POC and executed the end-to-end migration, achieving a **25-30% reduction** in infrastructure costs while maintaining system stability.
- Owned and managed Kafka Connect service, driving stabilization and performance enhancements. Automated on-boarding, alerting, and incident management, **reducing alerts by 90%**, improving system reliability.

EDUCATION

Indian Institute of Technology (IIT) BHU

Bachelor of Technology in Mining Engineering , CGPA: 8.70

Varanasi, India

July 2018 – May 2022

TECHNICAL SKILLS

Languages: Java, Golang, SQL, C++, Python, Scala

Data Frameworks and Tools: Spark, Airflow, ETL Pipelines, Kubernetes, Databricks, Kafka, Grafana, GCP, AWS

Frameworks & Databases: Spring Boot, Gin, Hadoop, Redis, MySQL, PostgreSQL

ACHIEVEMENTS

Panchayat Award: Recognized for leading Databricks to open-source notebook platform migration, AWS-to-GCP migration.

MMVY Scholar: Awarded a full tuition scholarship for academic excellence in JEE Main and Class 12.